

Generalized Linear Models with the `rstanarm` R Package

Ben Goodrich

April 08, 2019

Ex Ante Probability Density / Mass Function

A likelihood function is the same expression as a $P\{D,M\}F$ with 3 distinctions:

1. For the PDF or PMF, $f(x|\theta)$, we think of X as a random variable and θ as given, whereas we conceive of the likelihood function, $\mathcal{L}(\theta; x)$, to be a function of θ evaluated at the OBSERVED data, x
 - As a consequence, $\int_{-\infty}^{\infty} f(x|\theta) dx = 1$ or $\sum_{i: x_i \in \Omega} f(x_i|\theta) = 1$ while $\int_{-\infty}^{\infty} \int_{-\infty}^{\infty} \cdots \int_{-\infty}^{\infty} \mathcal{L}(\theta; x) d\theta_1 d\theta_2 \dots d\theta_K$ is positive but not 1
2. We often think of “the likelihood function” for N conditionally independent observations, so $\mathcal{L}(\theta; \mathbf{x}) = \prod_{n=1}^N \mathcal{L}(\theta; x_n)$
3. By “the likelihood function”, we often really mean the natural logarithm thereof: $\ell(\theta; \mathbf{x}) = \ln \mathcal{L}(\theta; \mathbf{x}) = \sum_{n=1}^N \ln \mathcal{L}(\theta; x_n)$
 - Fisher was concerned with the distribution of the mode across datasets

Hamiltonian Monte Carlo

- Instead of simply drawing from the posterior distribution whose PDF is $f(\boldsymbol{\theta}|\mathbf{y}) \propto f(\boldsymbol{\theta}) L(\boldsymbol{\theta}; \mathbf{y})$ Stan augments the “position” variables $\boldsymbol{\theta}$ with an equivalent number of “momentum” variables $\boldsymbol{\phi}$ and draws from

$$f(\boldsymbol{\theta}|\mathbf{y}) \propto \int_{-\infty}^{\infty} \cdots \int_{-\infty}^{\infty} \left[\prod_{k=1}^K \frac{1}{\sigma_k \sqrt{2\pi}} e^{-\frac{1}{2} \left(\frac{\phi_k}{\sigma_k} \right)^2} \right] f(\boldsymbol{\theta}) L(\boldsymbol{\theta}; \mathbf{y}) d\phi_1 \cdots d\phi_K$$

- Since the likelihood is NOT a function of ϕ_k , the posterior distribution of ϕ_k is the same as its prior, which is normal with a “tuned” standard deviation. So, at the s -th MCMC iteration, we just draw each $\tilde{\phi}_k$ from its normal distribution.
- Using physics, the realizations of each $\tilde{\phi}_k$ at iteration s “push” $\boldsymbol{\theta}$ from iteration $s - 1$ through the parameter space whose topology is defined by the negated log-kernel of the posterior distribution: $-\ln f(\boldsymbol{\theta}) - \ln L(\boldsymbol{\theta}; \mathbf{y})$

No U-Turn Sampling (NUTS)

- The location of θ moving according to Hamiltonian physics at any instant would be a valid draw from the posterior distribution
- But (in the absence of friction) θ moves indefinitely so when do you stop?
- Can simulate Hamiltonian dynamics “forward” and “backward” in “time”
- [Hoffman and Gelman \(2014\)](#) proposed stopping the forward dynamics and the backward dynamics start to get closer together. Hence, the name No U-Turn Sampling.
- After the U-Turn, one footprint is selected with probability proportional to the posterior kernel to be the realization of θ on iteration s and the process repeats itself
- NUTS discretizes a continuous-time Hamiltonian process in order to solve a system of Ordinary Differential Equations (ODEs), which requires a stepsize that is also tuned during the warmup phase

What is Stan?

- Includes a high-level [probabilistic programming language](#)
- Includes a translator of high-level Stan syntax to somewhat low-level C++
- Includes new (and old) gradient-based algorithms for statistical inference, such as NUTS
- Includes a matrix and scalar math library that supports autodifferentiation
- Includes interfaces from R and other high-level software
- Includes R packages with pre-written Stan programs
- Includes (not Stan specific) post-estimation R functions
- Includes a large community of users and many developers

Using Stan via R

1. Write the program in a (text) .stan file w/ R-like syntax that ultimately defines a posterior log-kernel. We will not do this until May. Stan's parser, `rstan::stanc`, does two things
 - checks that program is syntactically valid and tells you if not
 - writes a conceptually equivalent C++ source file to disk
 2. C++ compiler creates a binary file from the C++ source
 3. Execute the binary from R (can be concurrent with 2)
 4. Analyze the resulting samples from the posterior
 - Posterior predictive checks
 - Model comparison
 - Decision
- For the first several weeks, we are just going to focus on writing small functions in the Stan language

Sampling Distribution of OLS vs. Posterior Kernel

```
functions { /* saved as OLS_rng.stan*/  
  matrix OLS_rng(int S, real alpha, real beta,  
                 real sigma, vector x) {  
    matrix[S, 3] out; int N = rows(x);  
    vector[N] x_ = x - mean(x);  
    vector[N] mu = alpha + beta * x_;  
    real SSX = sum(square(x_));  
    for (s in 1:S) {  
      vector[N] y; vector[N] e;  
      for (n in 1:N)  
        y[n] = mu[n] + normal_rng(0, sigma);  
      out[s, 1] = mean(y);  
      out[s, 2] = sum(y .* x_) / SSX;  
      e = y - (out[s, 1] + out[s, 2] * x_);  
      out[s, 3] = sum(square(e)) / (N - 2);  
    }  
    return out;  
  }  
}
```


```
functions { /* saved as lm_kernel.stan*/  
  real lm_kernel(real alpha, real beta, real tau,  
                 vector y, vector x) {  
    int N = rows(x);  
    vector[N] x_ = x - mean(x);  
    vector[N] mu = alpha + beta * x_;  
  
    real sigma = inv_sqrt(tau);  
    //      ^^^ inv_sqrt(tau) = 1 / sqrt(tau)  
    // alpha and beta have improper priors ...  
    // ... so they add nothing to the log-kernel  
    return -log(tau) // Jeffreys prior on tau  
           + normal_lpdf(y | mu, sigma);  
    //      ^^^ log-likelihood of parameters  
  }  
}
```

Normal Distribution of the True Test Statistic

```
rstan::expose_stan_functions("OLS_rng.stan"); x <- lfactorial(0:16); alpha <- 0  
beta <- 0.5; sigma <- 10; OLS <- OLS_rng(S = 10 ^ 5, alpha, beta, sigma, x); colMeans(OLS)
```


```
## [1] -0.002718179 0.500792448 100.012222281
```

```
se <- sqrt(sigma ^ 2 / sum( (x - mean(x)) ^ 2 )) # true standard error of estimated slope  
plot(ecdf((OLS[ , 2] - beta) / se), main = "", xlab = "", ylab = "CDF of Test Statistic")  
curve(pnorm(z), from = -5, to = 5, lty = 2, col = 2, add = TRUE, xname = "z")  
legend("topleft", legend = c("Simulated", "Phi"), col = 1:2, lty = 1:2, bg = "lightgrey")
```



Student t Distribution of Estimated Test Statistic

```
se_hat <- sqrt(OLS[ , 3] / sum( (x - mean(x)) ^ 2 )) # estimated standard error of estimate
plot(ecdf((OLS[ , 2] - beta) / se_hat), main = "", xlab = "", ylab = "CDF of Test Statistic")
curve(pt(z, df = 17 - 2), from = -5, to = 5, lty = 2, col = 2, add = TRUE, xname = "z")
curve(pnorm(z), from = -5, to = 5, lty = 3, col = 3, add = TRUE, xname = "z")
legend("topleft", legend = c("Simulated", "True", "Phi"), col = 1:3, lty = 1:3, bg = "grey")
```



Power of the Test that $\beta = 0$ against $\beta > 0$

```
round(x, digits = 4)
```

```
## [1] 0.0000 0.0000 0.6931 1.7918 3.1781 4.7875 6.5793 8.5252 10.6046 12.8018  
## [11] 15.1044 17.5023 19.9872 22.5522 25.1912 27.8993 30.6719
```

```
mean( (OLS[ , 2] - 0) / se_hat > qt(0.95, df = 17 - 2) )
```

```
## [1] 0.62395
```

In other words, for THESE 17 values of x , we EXPECT (over Y) to reject the false null hypothesis that $\beta = 0$ in favor of the alternative hypothesis that $\beta > 0$ at the 5% level with probability 0.624 when the true value of β is $\frac{1}{2}$.

- What good is this PRE-DATA (on y_1, y_2, \dots, y_{17}) statement?
- But in this case the posterior distribution is the same as the estimated sampling distribution of the OLS estimator across datasets

Data on Diamonds

```
data("diamonds", package = "ggplot2")
str(diamonds)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    53940 obs. of  10 variables:
## $ carat   : num  0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
## $ cut     : Ord.factor w/ 5 levels "Fair"<"Good"<...: 5 4 2 4 2 3 3 3 1 3 ...
## $ color   : Ord.factor w/ 7 levels "D"<"E"<"F"<"G"<...: 2 2 2 6 7 7 6 5 2 5 ...
## $ clarity: Ord.factor w/ 8 levels "I1"<"SI2"<"SI1"<...: 2 3 5 4 2 6 7 3 4 5 ...
## $ depth   : num   61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
## $ table   : num    55 61 65 58 58 57 57 55 61 61 ...
## $ price    : int   326 326 327 334 335 336 336 337 337 338 ...
## $ x       : num    3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
## $ y       : num    3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
## $ z       : num    2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...
```

Generative Models

- In order to draw from the prior predictive distribution, you have to have a model that you can simulate from
- Without a generative model, you cannot update your beliefs with Bayes Rule

Concept	Known	Unknowable
Exogenous	sizes, predictors, prior modes / medians / etc.	parameters
Endogenous	outcomes	intermediates, predictions, utility

- **Endogenous Known:** Log of diamond prices
- **Exogenous Knowns:** N , physical characteristics of diamonds
- **Exogenous Unknowables :** α, β, R^2
- **Endogenous Unknowable :** (counterfactual?) predicted prices and functions thereof

Do This Once on Each Computer You Use

- R comes with a terrible default coding for ordered factors in regressions known as “Helmert” contrasts
- Execute this once to change them to “treatment” contrasts, which is the conventional coding in the social sciences with dummy variables relative to a baseline category

```
cat('options(contrasts = c(unordered = "contr.treatment", ordered = "contr.treatment"))',  
    file = "~/.Rprofile", sep = "\n", append = TRUE)
```

- Without this, you will get a weird rotation of the coefficients on the `cut` and `clarity` dummy variables

The stan_lm Function

```
library(rstanarm); options(mc.cores = parallel::detectCores())
post <- stan_lm(log(price) ~ carat * (log(x) + log(y) + log(z)) + cut + color + clarity, data = diamonds,
  prior = R2(location = 0.8, what = "mode"), subset = z > 0, adapt_delta = 0.95)
```

```
str(as.data.frame(post), vec.len = 3, digits.d = 2)
```

```
## 'data.frame':    4000 obs. of  28 variables:
## $ (Intercept)   : num  0.73 0.73 0.72 0.72 ...
## $ carat         : num  7.4 7.5 7.3 7.3 ...
## $ log(x)        : num  4.7 4.7 4.5 4.5 ...
## $ log(y)        : num  -2.6 -2.6 -2.5 -2.4 ...
## $ log(z)        : num  0.92 0.9 1.01 0.93 ...
## $ cutGood       : num  0.083 0.087 0.088 0.079 ...
## $ cutVery Good  : num  0.12 0.12 0.13 0.12 ...
## $ cutPremium    : num  0.13 0.14 0.14 0.13 ...
## $ cutIdeal      : num  0.16 0.17 0.17 0.16 ...
## $ colorE        : num  -0.057 -0.052 -0.053 -0.056 ...
## $ colorF        : num  -0.096 -0.096 -0.095 -0.096 ...
## $ colorG        : num  -0.16 -0.16 -0.16 -0.16 ...
## $ colorH        : num  -0.26 -0.26 -0.26 -0.26 ...
```

```
## $ colorI       : num  -0.37 -0.38 -0.37 -0.37 ...
## $ colorJ       : num  -0.51 -0.51 -0.51 -0.51 ...
## $ claritySI2   : num  0.41 0.41 0.42 0.42 ...
## $ claritySI1   : num  0.58 0.58 0.59 0.58 ...
## $ clarityVS2   : num  0.73 0.73 0.73 0.73 ...
## $ clarityVS1   : num  0.8 0.8 0.8 0.8 ...
## $ clarityVVS2  : num  0.93 0.93 0.94 0.93 ...
## $ clarityVVS1  : num  1 1 1 1 ...
## $ clarityIF    : num  1.1 1.1 1.1 1.1 ...
## $ carat:log(x) : num  -4.1 -4.1 -3.9 -3.9 ...
## $ carat:log(y) : num  2 1.9 1.9 1.9 ...
## $ carat:log(z) : num  -1.1 -1 -1.1 -1.1 ...
## $ sigma        : num  0.13 0.13 0.13 0.13 ...
## $ log-fit_ratio: num  -6.5e-04 -6.6e-05 2.7e-04 2.0e-04 ...
## $ R2           : num  0.98 0.98 0.98 0.98 ...
```

Typical Output

```
print(post, digits = 4)
```

```
...  
##           Median MAD_SD  
## (Intercept)   0.7664  0.0355  
## carat         7.4186  0.0724  
## log(x)        4.5230  0.0804  
## log(y)       -2.5166  0.0721  
## log(z)        0.9599  0.0424  
## cutGood       0.0852  0.0038  
## cutVery Good  0.1223  0.0037  
## cutPremium    0.1353  0.0036  
## cutIdeal      0.1665  0.0036  
## colorE       -0.0551  0.0020  
## colorF       -0.0961  0.0020  
## colorG       -0.1628  0.0019  
## colorH       -0.2572  0.0021  
## colorI       -0.3750  0.0024  
## colorJ       -0.5116  0.0029  
## claritySI2    0.4166  0.0050  
## claritySI1    0.5821  0.0049
```

```
## clarityVS2    0.7290  0.0050  
## clarityVS1    0.8001  0.0050  
## clarityVVS2   0.9309  0.0051  
## clarityVVS1   1.0022  0.0053  
## clarityIF     1.0974  0.0059  
## carat:log(x)  -3.9631  0.0661  
## carat:log(y)   1.9113  0.0568  
## carat:log(z)  -1.1213  0.0443  
## sigma         0.1257  0.0004  
## log-fit_ratio  0.0000  0.0005  
## R2            0.9846  0.0001  
##  
## Sample avg. posterior predictive distribution  
##           Median MAD_SD  
## mean_PPD  7.7864  0.0007  
##  
## -----  
## For info on the priors used see help('prior_s  
...
```

More Detailed Output

```
print(summary(post), digits = 3) # shows estimated effective sample sizes at the bottom
```

```
##
## Model Info:
##
## function:      stan_lm
## family:        gaussian [identity]
## formula:       log(price) ~ carat * (log(x) + log(y) + log(z)) + cut + color +
##               clarity
## algorithm:     sampling
## priors:         see help('prior_summary')
## sample:        4000 (posterior sample size)
## observations:  53920
## predictors:    25
##
## Estimates:
##
```

	mean	sd	2.5%	25%	50%	75%	97.5%
## (Intercept)	0.767	0.035	0.699	0.743	0.766	0.791	0.836
## carat	7.419	0.074	7.276	7.370	7.419	7.468	7.562
## log(x)	4.523	0.080	4.368	4.469	4.523	4.578	4.682
## log(y)	-2.517	0.075	-2.660	-2.565	-2.517	-2.468	-2.371
## log(z)	0.960	0.042	0.878	0.931	0.960	0.989	1.043
## cutGood	0.085	0.004	0.078	0.083	0.085	0.088	0.093
## cutVery Good	0.122	0.004	0.115	0.120	0.122	0.125	0.130
## cutPremium	0.135	0.004	0.128	0.133	0.135	0.138	0.142
## cutIdeal	0.167	0.004	0.160	0.164	0.167	0.169	0.173
## colorE	-0.055	0.002	-0.059	-0.056	-0.055	-0.054	-0.051
## colorF	-0.096	0.002	-0.100	-0.098	-0.096	-0.095	-0.092
## colorG	-0.163	0.002	-0.167	-0.164	-0.163	-0.162	-0.159
## colorH	-0.257	0.002	-0.261	-0.259	-0.257	-0.256	-0.253
## colorI	-0.375	0.002	-0.379	-0.377	-0.375	-0.373	-0.370
## colorJ	-0.512	0.003	-0.517	-0.514	-0.512	-0.510	-0.506
## claritySI2	0.417	0.005	0.407	0.413	0.417	0.420	0.426

Credible Intervals


```
round(posterior_interval(post, prob = 0.8),  
      digits = 2)
```

##	10%	90%
## (Intercept)	0.72	0.81
## carat	7.32	7.51
## log(x)	4.42	4.62
## log(y)	-2.61	-2.42
## log(z)	0.90	1.01
## cutGood	0.08	0.09
## cutVery Good	0.12	0.13
## cutPremium	0.13	0.14
## cutIdeal	0.16	0.17
## colorE	-0.06	-0.05
## colorF	-0.10	-0.09
## colorG	-0.17	-0.16
## colorH	-0.26	-0.25

## colorI	-0.38	-0.37
## colorJ	-0.52	-0.51
## claritySI2	0.41	0.42
## claritySI1	0.58	0.59
## clarityVS2	0.72	0.74
## clarityVS1	0.79	0.81
## clarityVVS2	0.92	0.94
## clarityVVS1	1.00	1.01
## clarityIF	1.09	1.10
## carat:log(x)	-4.05	-3.88
## carat:log(y)	1.84	1.98
## carat:log(z)	-1.18	-1.06
## sigma	0.13	0.13
## log-fit_ratio	0.00	0.00
## R2	0.98	0.98

What Is the Effect of an Increase in Carat?

```
PPD_0 <- exp(posterior_predict(post, draws = 500))  
df <- diamonds[diamonds$z > 0, ]; df$carat <- df$carat + 0.2  
PPD_1 <- exp(posterior_predict(post, newdata = df, draws = 500))  
plot(ecdf((PPD_1 - PPD_0) / PPD_0), main = "", xlim = c(0, 2), xlab = "Ratio", ylab = "CDF")
```



Why NUTS Is Better than Other MCMC Samplers

- With Stan, it is almost always the case that things either go well or you get warning messages saying things did not go well
- Because Stan uses gradients, it scales well as models get more complex
- The first-order autocorrelation tends to be negative so you can get greater effective sample sizes (for mean estimates) than nominal sample size

```
round(bayesplot::neff_ratio(post), digits = 2)
```

##	(Intercept)	carat	log(x)	log(y)	log(z)	cutGood
##	0.44	0.78	0.91	1.03	0.83	1.08
##	cutVery Good	cutPremium	cutIdeal	colorE	colorF	colorG
##	1.11	1.14	1.11	1.08	1.02	1.07
##	colorH	colorI	colorJ	claritySI2	claritySI1	clarityVS2
##	0.69	1.11	0.89	1.05	1.04	1.03
##	clarityVS1	clarityVVS2	clarityVVS1	clarityIF	carat:log(x)	carat:log(y)
##	1.01	0.96	1.03	1.01	1.04	1.02
##	carat:log(z)	sigma	log-fit_ratio	R2		
##	0.57	0.43	1.08	0.42		

Divergent Transitions

- NUTS only uses first derivatives
- First order approximations to Hamiltonian physics are fine for if either the second derivatives are constant or the discrete step size is sufficiently small
- When the second derivatives are very not constant across Θ , Stan can (easily) mis-tune to a step size that is not sufficiently small and θ_k gets pushed to $\pm\infty$
- When this happens there will be a warning message, suggesting to increase `adapt_delta`
- When `adapt_delta` is closer to 1, Stan will tend to take smaller steps
- Unfortunately, even as `adapt_delta` $\lim 1$, there may be no sufficiently small step size and you need to try to reparameterize your model

Exceeding Maximum Treedepth

- When the step size is small, NUTS needs many (small) steps to cross the “typical” subset of Θ and hit the U-turn point
- Sometimes, NUTS has not U-turned when it reaches its limit of 10 steps (by default)
- When this happens there will be a warning message, suggesting to increase `max_treedepth`
- There is always a sufficiently high value of `max_treedepth` to allow NUTS to reach the U-turn point, but increasing `max_treedepth` by 1 approximately doubles the wall time to obtain S draws

The `stan_glm` Function in `rstanarm`


- All examples from the reading (plus more) are available at <https://github.com/avehtari/RAOS-Examples>

```
library(rstanarm); options(mc.cores = parallel::detectCores());
data(kidiq, package = "rstanarm")
fit_4 <- stan_glm(kid_score ~ mom_hs * mom_iq, data = kidiq, family = gaussian(),
                 prior_intercept = normal(location = 100, scale = 15, autoscale = FALSE),
                 prior = normal(autoscale = FALSE), QR = TRUE)
```

fit_4


```
...
##               Median MAD_SD
## (Intercept)  -11.1    14.0
## mom_hs       51.1    15.5
## mom_iq        1.0     0.1
## mom_hs:mom_iq -0.5     0.2
## sigma       18.0     0.6
##
## Sample avg. posterior predictive distribution of y:
##               Median MAD_SD
## mean_PPD 86.8     1.2
##
```

Plot at the Posterior Median Estimates



Correct Plot


```
pp_check(fit_4, plotfun = "loo_intervals", order = "median")
```



What Did the Priors Imply?

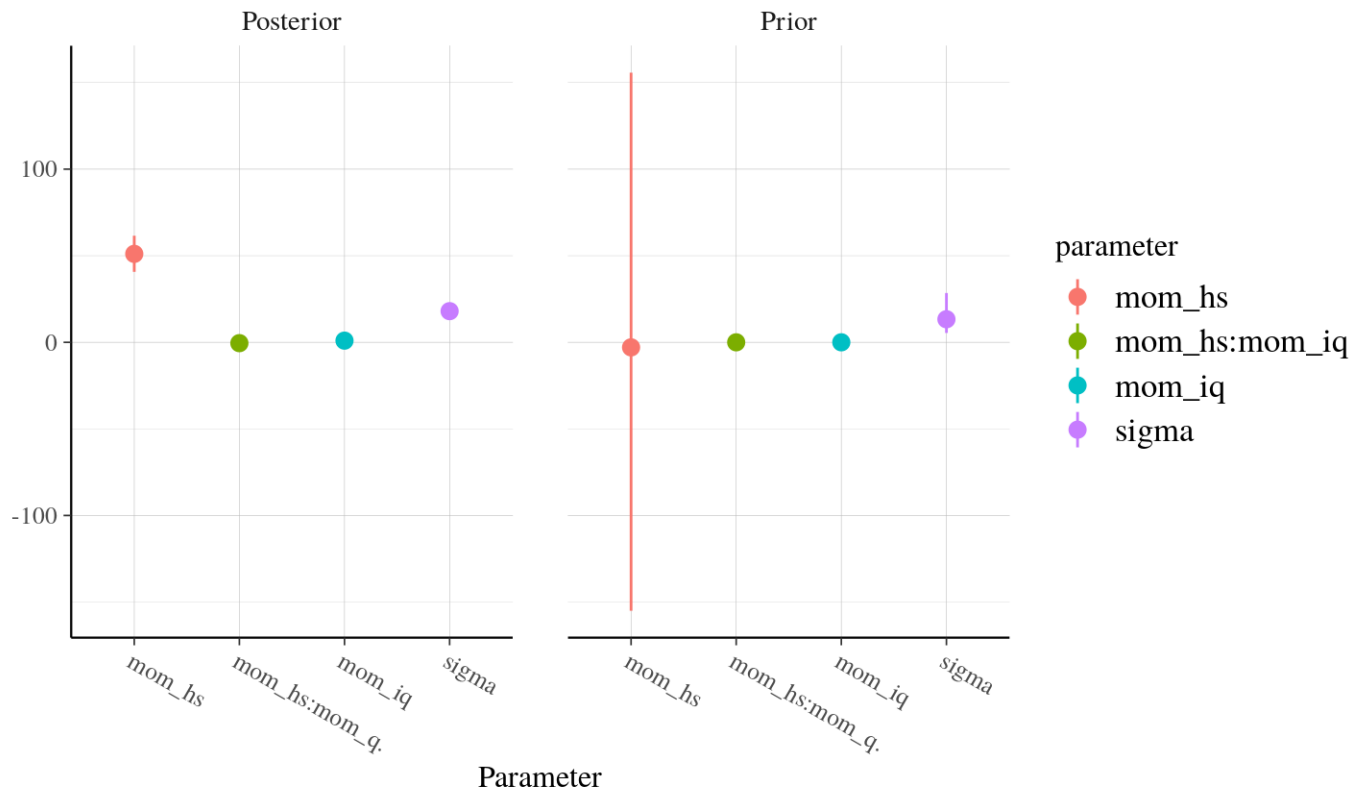
```
prior_4 <- update(fit_4, prior_PD = TRUE) # calls stan_glm() again with new arguments
```

```
plot(prior_4)
```



Posterior vs. Prior

```
posterior_vs_prior(fit_4, prob = 0.5, regex_pars = "^[^(]") # excludes (Intercept)
```



Logit Models with Normal Priors

```
functions { /* saved as logit_PPD_rng.stan */
  matrix
  logit_PPD_rng(int S, vector x, real a_loc,
                real a_scale, real b_loc,
                real b_scale) {
    int N = rows(x); matrix[S, N] y_tilde;
    vector[N] x_ = x - mean(x);
    for (s in 1:S) {
      real alpha = normal_rng(a_loc, a_scale);
      real beta = normal_rng(b_loc, b_scale);
      vector[N] eta = alpha + beta * x_;
      for (n in 1:N) {
        real utils = eta[n] + logistic_rng(0, 1);
        y_tilde[s, n] = utils > 0;
      }
    }
    return y_tilde;
  }
}
```

```
functions { /* saved as logit_kernel.stan */
  real
  logit_kernel(real alpha, real beta,
               real a_loc, real a_scale,
               real b_loc, real b_scale,
               int[] y, vector x) {
    int N = rows(x); vector[N] x_ = x - mean(x);

    real p = normal_lpdf(alpha | a_loc, a_scale);
    real q = normal_lpdf(beta | b_loc, b_scale);
    vector[N] eta = alpha + beta * x_;


    return p + q + // priors & log-likelihood
      bernoulli_logit_lpmf(y | eta);
  }
}
```

What Does `bernoulli_logit_lpmf` Do?

- What is the logarithm of the Bernoulli PMF?
- The Bernoulli PMF is $\Pr(y \mid \mu) = \mu^y(1 - \mu)^{1-y}$, so its logarithm is $\ell(\mu; y) = y \ln \mu + (1 - y) \ln(1 - \mu)$
- What is the conditional distribution of utility in a logit model?
- What is an expression for probability that $Y = 1$ if utility is logistic with expectation η and scale 1?
- $\Pr(Y = 1) = \Pr(\eta + \epsilon > 0) = 1 - F(0 \mid \eta, 1) = F(\eta \mid 0, 1) = \frac{1}{1 + e^{-\eta}}$
- Combining these, we get $\ell(\mu; y) = -y \ln(1 + e^{-\eta}) + (1 - y)(-\eta - \ln(1 + e^{-\eta}))$
- Can save time by only calculating $\ln(1 + e^{-\eta})$ once and can preserve numerical precision by using `log1p_exp(-eta)`

Checking the Prior Predictive Distribution

```
rstan::expose_stan_functions("logit_PPD_rng.stan"); N <- 100; x <- 1:N  
y_tilde <- logit_PPD_rng(S = 10000, x, a_loc = 0, a_scale = 4, b_loc = 2, b_scale = 4)  
ggplot(data.frame(x, y_tilde = colMeans(y_tilde))) + geom_smooth(aes(x = x, y = y_tilde))
```



A Logit Model for Romney vs Obama in 2012

```
poll <- readRDS("GooglePoll.rds") # WantToWin is coded as 1 for Romney and 0 for Obama
library(dplyr)
collapsed <- filter(poll, !is.na(WantToWin)) %>%
  group_by(Region, Gender, Urban_Density, Age, Income) %>%
  summarize(Romney = sum(grepl("Romney", WantToWin)), Obama = n() - Romney) %>%
  na.omit
```

```
post <- stan_glm(cbind(Romney, Obama) ~ ., data = collapsed, family = binomial(link = "logit"), QR = TRUE)
```

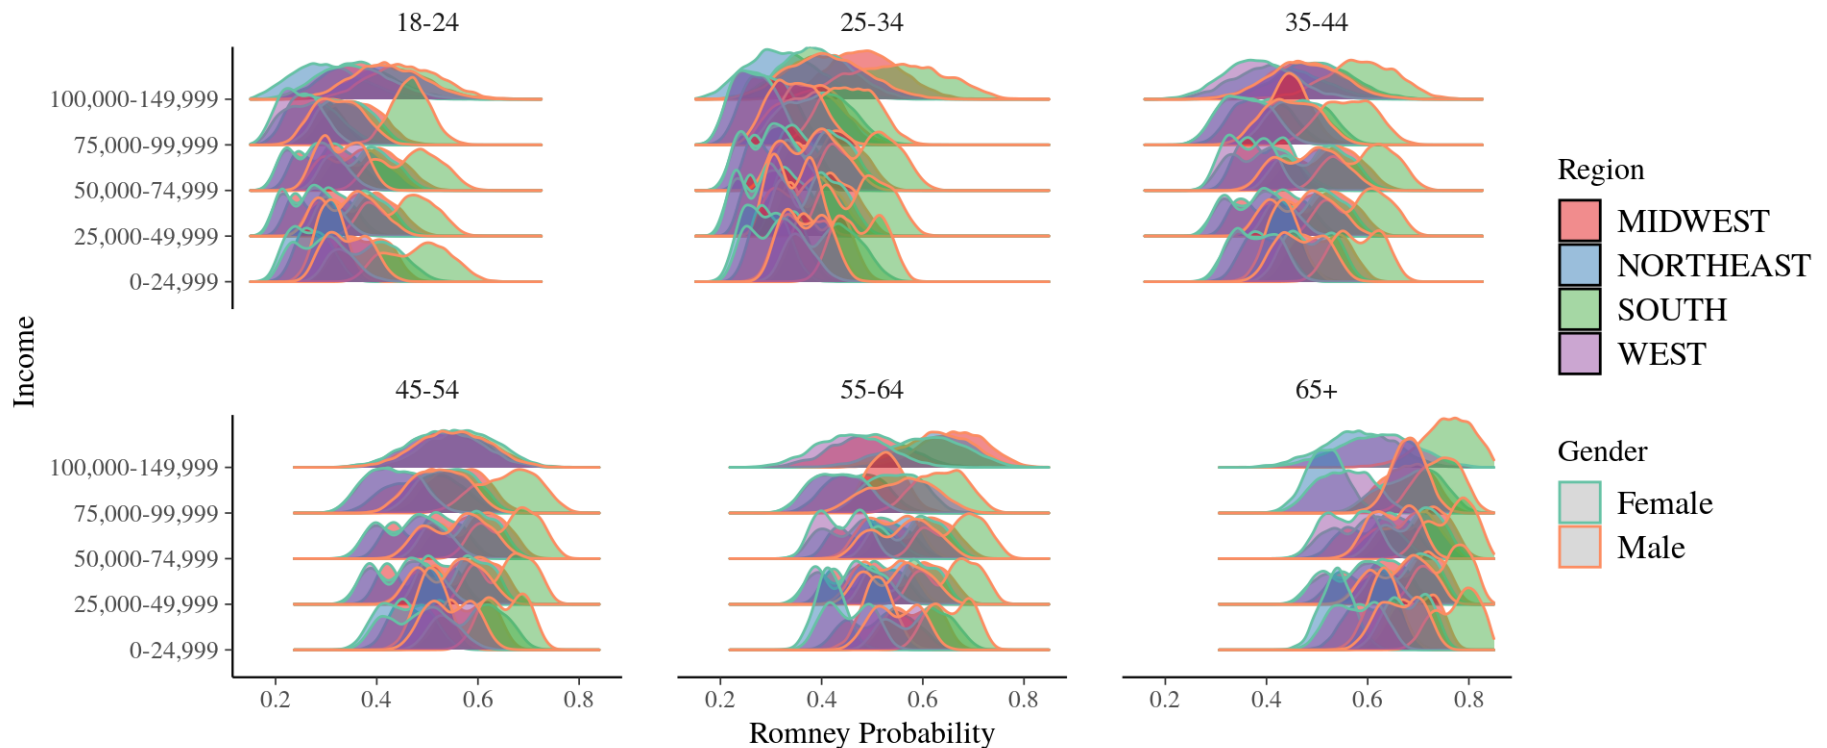
```
print(post, digits = 2)
```

```
...
##               Median MAD_SD
## (Intercept)    -0.53   0.14
## RegionNORTHEAST -0.14   0.09
## RegionSOUTH      0.31   0.07
## RegionWEST       -0.14   0.07
## GenderMale       0.38   0.06
## Urban_DensitySuburban -0.21 0.09
## Urban_DensityUrban  -0.52 0.09
```

```
## Age25-34          0.11  0.10
## Age35-44          0.54  0.10
## Age45-54          0.84  0.09
## Age55-64          0.85  0.09
## Age65+            1.35  0.11
## Income25,000-49,999 -0.12 0.08
## Income50,000-74,999 -0.07 0.09
## Income75,000-99,999 -0.09 0.14
## Income100,000-149,999 0.17 0.29
## Income150,000+      0.84 1.01
...
```

Posterior Distribution of Conditional Expectations

```
library(tidybayes); filter(collapsed, Income != "150,000+") %>% add_fitted_draws(post) %>%  
  ggplot(aes(x = .value, y = Income, fill = Region, color = Gender)) + xlim(0.15, 0.85) +  
  scale_fill_brewer(palette = "Set1") + scale_color_brewer(palette = "Set2") +  
  ggribes::geom_density_ridges(alpha = .5) + facet_wrap(~ Age) + xlab("Romney Probability")
```



Generalized Linear Models, in General

- BEFORE you see the outcome data Y has some probability distribution, which for Frequentists must be in the exponential family but that includes a lot of familiar probability distributions
- That probability distribution has a location parameter, which for Frequentists must be the conditional expectation given the predictors
- There is a one-to-one “link” function mapping from this conditional expectation, $\mu(x)$ to the “linear predictor”, $\eta(x)$, although it is often more natural to think about the “inverse link” function that maps from the “linear predictor”, $\eta(x)$ to the conditional expectation, $\mu(x)$
- The “linear predictor”, $\eta(x)$, must be a linear function of the PARAMETERS for Frequentists, although this is not necessary for Bayesians

Generalized Linear Model Examples This Week

Model	<i>Ex ante</i> Outcome Distribution	Inverse Link
Linear	<code>normal_lpdf(y mu, sigma)</code>	Identity: $\mu = \eta$
Logit	<code>bernoulli_lpmf(y mu)</code>	Inverse logit: $\mu = \frac{1}{1+e^{-\eta}}$
Probit	<code>bernoulli_lpmf(y mu)</code>	Std. Normal: $\mu = \Phi(\eta)$
Poisson	<code>poisson_lpmf(y mu)</code>	Antilog: $\mu = e^{\eta}$
Neg Binomial*	<code>neg_binomial_2_lpmf(y mu, phi)</code>	Antilog: $\mu = e^{\eta}$

-
- Not actually a Generalized Linear Model in the Frequentist sense
 - In each case, $\eta_i = \alpha + \sum_{k=1}^K x_i \beta_k$
 - The normal priors on α and each β_k could be changed to (an)other family / families

Poisson Models with Normal Priors

```
functions { /* named poisson_PPD_rng.stan */  
  matrix poisson_PPD_rng(int S, vector x,  
                          real a_loc,  
                          real a_scale,  
                          real b_loc,  
                          real b_scale) {  
    int N = rows(x); matrix[S, N] y_tilde;  
    vector[N] x_ = x - mean(x);  
    for (s in 1:S) {  
      real alpha = normal_rng(a_loc, a_scale);  
      real beta = normal_rng(b_loc, b_scale);  
      vector[N] eta = alpha + beta * x_;  
      vector[N] mu = exp(eta);  
      for (n in 1:N)  
        y_tilde[s, n] = poisson_rng(mu[n]);  
    }  
    return y_tilde;  
  }  
}
```

```
functions { /* named poisson_kernel.stan */  
  real poisson_kernel(real alpha, real beta,  
                      real a_loc,  
                      real a_scale,  
                      real b_loc,  
                      real b_scale,  
                      int[] y, vector x) {  
    int N = rows(x); vector[N] x_ = x - mean(x);  
  
    real p = normal_lpdf(alpha | a_loc, a_scale);  
    real q = normal_lpdf(beta | b_loc, b_scale);  
    vector[N] eta = alpha + beta * x_;  
  
    return p + q + // priors & log-likelihood  
      poisson_log_lpmf(y | eta);  
  }  
}
```

Negative Binomial Models with Normal Priors

```
functions { /* saved as nb_PPD_rng.stan */
  matrix
  nb_PPD_rng(int S, vector x, real a_loc,
             real a_scale, real b_loc,
             real b_scale, real rate) {
    int N = rows(x); matrix[S, N] y_tilde;
    vector[N] x_ = x - mean(x); for (s in 1:S) {
      real alpha = normal_rng(a_loc, a_scale);
      real beta = normal_rng(b_loc, b_scale);
      real phi = 1 / exponential_rng(rate);
      for (n in 1:N) {
        real z = exp(alpha + beta * x_[n]) / phi;
        real lambda = gamma_rng(z, z);
        y_tilde[s, n] = poisson_rng(lambda);
      }
    }
    return y_tilde;
  }
}
```

```
functions { /* saved as nb_kernel.stan */
  real nb_PPD_rng(real alpha, real beta,
                 real inv_phi, real a_loc,
                 real a_scale, real b_loc,
                 real b_scale, real rate,
                 vector x, int[] y) {
    int N = rows(x); vector[N] x_ = x - mean(x);
    real p = normal_lpdf(alpha | a_loc, a_scale);
    real q = normal_lpdf(beta | b_loc, b_scale);
    real r = exponential_lpdf(inv_phi | rate);
    vector[N] eta = alpha + beta * x_;

    return p + q + r + // priors & loglikelihood
      neg_binomial_2_log_lpmf(y | eta,
                             inv(inv_phi));
  }
}
```

Count Models with `stan_glm{.nb}`

```
post <- stan_glm.nb(y ~ I(roach1 / 100) + treatment + senior, data = roaches,  
  offset = log(exposure2), QR = TRUE)
```

post

```
...  
##  
##           Median MAD_SD  
## (Intercept)      2.8    0.2  
## I(roach1/100)     1.3    0.2  
## treatment        -0.8    0.3  
## senior           -0.3    0.3  
## reciprocal_dispersion 0.3    0.0  
##  
## Sample avg. posterior predictive distribution of y:  
##           Median MAD_SD  
## mean_PPD 49.1    28.7  
##  
## -----  
## For info on the priors used see help('prior_summary.stanreg').  
...
```

Estimating Treatment Effects

```
df <- roaches; df$treatment <- 0  
Y_0 <- posterior_linpred(post, newdata = df, offset = log(df$exposure2), transform = TRUE)  
df$treatment <- 1  
Y_1 <- posterior_linpred(post, newdata = df, offset = log(df$exposure2), transform = TRUE)  
plot(ecdf((Y_1 - Y_0) / Y_0), pch = ".", xlab = "Relative treatment effect", main = "")
```

