

# What Stan Does

Ben Goodrich

February 21, 2019

# Obligatory Disclosure

- Ben is an employee of Columbia University, which has received several research grants to develop Stan and is currently [seeking](#) an additional \$10 million
- Ben is also a manager of GG Statistics LLC, which uses Stan for business purposes
- According to Columbia University [policy](#), any such employee who has any equity stake in, a title (such as officer or director) with, or is expected to earn at least \$5,000.00 per year from a private company is required to disclose these facts in presentations

# Review of Ancient MCMC Samplers

- Metropolis-Hastings (M-H)
  - Only requires user to specify numerator of Bayes Rule
  - But only 22% of proposals ideally get accepted to get relatively big jumps
  - Effective Sample Size ( $n_{eff}$ ) can be essentially zero
- Gibbs sampling (in general)
  - User has to work out all full-conditional distributions
  - Jumps always accepted but might not be very big
  - Effective Sample Size is low if the parameters are highly correlated
- What the BUGS family (WinBUGS, JAGS, OpenBUGS) of software does
  - User does not have to work out full-conditional distributions
  - Falls back to some other algorithm (such as M-H but not actually M-H) when kernel of  $k$ -th full-conditional distribution is not recognized

# Comparing Stan to Ancient MCMC Samplers

- Like M-H, only requires user to specify numerator of Bayes Rule
- Like M-H but unlike Gibbs sampling, proposals are joint
- Unlike M-H but like Gibbs sampling, proposals always accepted
- Unlike M-H but like Gibbs sampling, tuning of proposals is (often) not required
- Unlike both M-H and Gibbs sampling, the effective sample size is typically 25% to 125% of the nominal number of draws from the posterior distribution because  $\rho_1$  can be negative in  $n_{eff} = \frac{S}{1 + 2 \sum_{k=1}^{\infty} \rho_k}$
- Unlike both M-H and Gibbs sampling, Stan produces warning messages when things are not going swimmingly. Do not ignore these!
- Unlike BUGS, Stan does not permit discrete unknowns but even BUGS has difficulty drawing discrete unknowns with a sufficient amount of efficiency

# Hamiltonian Monte Carlo

- Instead of simply drawing from the posterior distribution whose PDF is  $f(\boldsymbol{\theta} | \mathbf{y} \dots) \propto f(\boldsymbol{\theta}) L(\boldsymbol{\theta}; \mathbf{y})$  Stan augments the “position” variables  $\boldsymbol{\theta}$  with an equivalent number of “momentum” variables  $\boldsymbol{\phi}$  and draws from

$$f(\boldsymbol{\theta} | \mathbf{y} \dots) \propto \int_{-\infty}^{\infty} \dots \int_{-\infty}^{\infty} \prod_{k=1}^K \frac{1}{\sigma_k \sqrt{2\pi}} e^{-\frac{1}{2} \left( \frac{\phi_k}{\sigma_k} \right)^2} f(\boldsymbol{\theta}) L(\boldsymbol{\theta}; \mathbf{y}) d\phi_1 \dots d\phi_K$$

- Since the likelihood is NOT a function of  $\phi_k$ , the posterior distribution of  $\phi_k$  is the same as its prior, which is normal with a “tuned” standard deviation. So, at the  $s$ -th MCMC iteration, we just draw each  $\tilde{\phi}_k$  from its normal distribution.
- Using physics, the realizations of each  $\tilde{\phi}_k$  at iteration  $s$  “push”  $\boldsymbol{\theta}$  from iteration  $s - 1$  through the parameter space whose topology is defined by the negated log-kernel of the posterior distribution:  $-\ln f(\boldsymbol{\theta}) - \ln L(\boldsymbol{\theta}; \mathbf{y})$
- See HMC.R demo on Canvas

# Demo of Hamiltonian Monte Carlo



# No U-Turn Sampling (NUTS)

- The location of  $\theta$  moving according to Hamiltonian physics at any instant would be a valid draw from the posterior distribution
- But (in the absence of friction)  $\theta$  moves indefinitely so when do you stop?
- [Hoffman and Gelman \(2014\)](#) proposed stopping when there is a “U-turn” in the sense the footprints turn around and start to head in the direction they just came from. Hence, the name No U-Turn Sampling.
- After the U-Turn, one footprint is selected with probability proportional to the posterior kernel to be the realization of  $\theta$  on iteration  $s$  and the process repeats itself
- NUTS discretizes a continuous-time Hamiltonian process in order to solve a system of Ordinary Differential Equations (ODEs), which requires a stepsize that is also tuned during the warmup phase

# What is Stan?

- Includes a high-level [probabilistic programming language](#)
- Includes a translator of high-level Stan syntax to somewhat low-level C++
- Includes new (and old) gradient-based algorithms for statistical inference, such as NUTS
- Includes a matrix and scalar math library that supports autodifferentiation
- Includes interfaces from R and other high-level software
- Includes R packages with pre-written Stan programs
- Includes (not Stan specific) post-estimation R functions
- Includes a large community of users and many developers



# What is Autodifferentiation?

- A language like C++ supports operator overloading of  $+$ ,  $-$ , etc. to do whatever
- In Stan,  $c = a / b$  computes  $c$  and both  $\frac{\partial c}{\partial a} = \frac{1}{b}$  and  $\frac{\partial c}{\partial b} = -\frac{a}{b^2}$
- Similarly,  $d = g(c)$  computes  $d$  and  $\frac{\partial d}{\partial c} = g'(c)$
- Evaluating the chain rule is tedious for a human but easy for a computer
- Autodifferentiation allows the human to write an arbitrary (differentiable) mathematical expression and the (C++) compiler generates the code to compute the derivative automatically, even with vector / matrix expressions
- This is more accurate and / or faster than symbolic differentiation or numerical differentiation
- For Stan's purposes, Stan does autodifferentiation faster than anything else; see <http://arxiv.org/abs/1509.07164>

# Using Stan via R

1. Write the program in a (text) .stan file w/ R-like syntax that ultimately defines a posterior log-kernel. We will not do this until April. Stan's parser, `rstan::stanc`, does two things
  - checks that program is syntactically valid & tells you if not
  - writes a conceptually equivalent C++ source file to disk
2. C++ compiler creates a binary file from the C++ source
3. Execute the binary from R (can be concurrent with 2)
4. Analyze the resulting samples from the posterior
  - Posterior predictive checks
  - Model comparison
  - Decision

# Drawing from a Bivariate Normal with NUTS

```
library(rstan)
xy <- stan("binormal.stan", refresh = 0)
xy

## Inference for Stan model: binormal.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##          mean se_mean   sd  2.5%  25%   50%   75% 97.5% n_eff Rhat
## x      0.03    0.03 1.02 -2.01 -0.66  0.05  0.72  1.98  1123 1.00
## y      0.02    0.03 1.01 -1.98 -0.64  0.03  0.70  1.99  1178 1.00
## lp__ -2.44    0.03 1.06 -5.22 -2.83 -2.11 -1.71 -1.45  1093 1.01
##
## Samples were drawn using NUTS(diag_e) at Thu Feb 21 15:22:04 2019.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

# Divergent Transitions

- NUTS only uses first derivatives
- First order approximations to Hamiltonian physics are fine for if either the second derivatives are constant or the discrete step size is sufficiently small
- When the second derivatives are very not constant across  $\Theta$ , Stan can (easily) mis-tune to a step size that is not sufficiently small and  $\theta_k$  gets pushed to  $\pm\infty$
- When this happens there will be a warning message, suggesting to increase `adapt_delta`
- When `adapt_delta` is closer to 1, Stan will tend to take smaller steps
- Unfortunately, even as `adapt_delta`  $\lim 1$ , there may be no sufficiently small step size and you need to try to reparameterize your model

# Exceeding Maximum Treedepth

- When the step size is small, NUTS needs many (small) steps to cross the “typical” subset of  $\Theta$  and hit the U-turn point
- Sometimes, NUTS has not U-turned when it reaches its limit of 10 steps (by default)
- When this happens there will be a warning message, suggesting to increase `max_treedepth`
- There is always a sufficiently high value of `max_treedepth` to allow NUTS to reach the U-turn point, but increasing `max_treedepth` by 1 approximately doubles the wall time to obtain  $S$  draws

# Low Bayesian Fraction of Missing Information

- When the tails of the posterior PDF are very light, NUTS can have difficulty moving through  $\Theta$  efficiently
- This will manifest itself in a low (and possibly unreliable) estimate of  $n_{eff}$
- When this happens there will be a warning message, saying that the Bayesian Fraction of Missing Information (BFMI) is low
- In this situation, there is not much you can do except increase  $S$  or preferably reparameterize your model to make it easier for NUTS

# Runtime Exceptions

- Sometimes you will get a “informational” (not error, not warning) message saying that some parameter that should be positive is zero or some parameter that should be finite is infinite
- This means that a 64bit computer could not represent the number accurately
- If it only happens a few times and only during the warmup phase, do not worry
- Otherwise, you might try to use functions that are more numerically stable, which is discussed throughout the Stan User Manual